

LAB-02

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
In [3]: df= pd.read_csv("C:/Users/91969/Downloads/train.csv")
df.head()
```

Out[3]:

	trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees	total
0	748.0	2.75	1.0	75.00	24	6.300	10
1	1187.0	3.43	1.0	105.00	24	13.200	14
2	730.0	3.12	1.0	71.25	0	26.625	9
3	671.0	5.63	3.0	90.00	0	9.750	9
4	329.0	2.09	1.0	45.00	12	13.200	7

```
In [3]: df.shape
```

Out[3]: (209673, 8)

```
In [4]: df.describe()
```

Out[4]:

	trip_duration	distance_traveled	num_of_passengers	fare	tip
count	209673.000000	209673.000000	209673.000000	209673.000000	209673.000000
mean	1173.181478	5.054431	1.292808	99.623431	13.030824
std	4775.653621	125.217419	0.930754	85.602702	20.367764
min	0.000000	0.020000	0.000000	0.000000	0.000000
25%	446.000000	1.950000	1.000000	52.500000	0.000000
50%	707.000000	3.200000	1.000000	75.000000	9.000000
75%	1098.000000	5.730000	1.000000	116.250000	20.000000
max	86387.000000	57283.910000	9.000000	4466.250000	2500.000000

In [5]: `df.info()`


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209673 entries, 0 to 209672
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   trip_duration         209673 non-null float64
1   distance_traveled     209673 non-null float64
2   num_of_passengers     209673 non-null float64
3   fare                  209673 non-null float64
4   tip                   209673 non-null int64  
5   miscellaneous_fees    209673 non-null float64
6   total_fare            209673 non-null float64
7   surge_applied         209673 non-null int64  
dtypes: float64(6), int64(2)
memory usage: 12.8 MB
```

In [6]: `df.isnull().sum()`

```
Out[6]: trip_duration         0
distance_traveled         0
num_of_passengers         0
fare                     0
tip                      0
miscellaneous_fees        0
total_fare                0
surge_applied             0
dtype: int64
```

Data Visualization

```
In [7]: plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)  
plt.show()
```



Train and Test Split

```
In [8]: X = df.drop(columns='total_fare')
X
```

Out[8]:

	trip_duration	distance_traveled	num_of_passengers	fare	tip	miscellaneous_fees
0	748.0	2.75	1.0	75.00	24	6.300
1	1187.0	3.43	1.0	105.00	24	13.200
2	730.0	3.12	1.0	71.25	0	26.625
3	671.0	5.63	3.0	90.00	0	9.750
4	329.0	2.09	1.0	45.00	12	13.200
...
209668	1617.0	8.42	1.0	150.00	47	5.800
209669	438.0	1.29	1.0	48.75	12	34.575
209670	571.0	2.82	1.0	63.75	0	6.000
209671	491.0	2.16	1.0	56.25	0	13.500
209672	3614.0	33.72	1.0	337.50	0	2.250

209673 rows × 7 columns



```
In [9]: y = df['total_fare']
y
```

```
Out[9]: 0      105.300
1      142.200
2       97.875
3       99.750
4       70.200
...
209668  202.800
209669   95.325
209670   69.750
209671   69.750
209672  339.750
Name: total_fare, Length: 209673, dtype: float64
```

```
In [10]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
```

Feature Scaling

```
In [11]: STD = StandardScaler()
X_train_SD = STD.fit_transform(X_train)
X_test_SD = STD.transform(X_test)
```

Model Training

```
In [12]: LR = LinearRegression()  
LR.fit(X_train_SD, y_train)
```

```
Out[12]: LinearRegression()
```

```
In [13]: y_hat = LR.predict(X_test_SD)
```

Model Evaluation

```
In [14]: print("mse: ", mean_squared_error(y_hat, y_test))
```

```
mse:  3.473722029833163e-27
```


```
In [15]: regressor = LinearRegression()  
regressor.fit(X_train, y_train)  
accuracy = regressor.score(X_test, y_test)  
print("Accuracy:",(int(round(accuracy*100))))
```

```
Accuracy: 100
```

```
In [16]: print("mae: ", mean_absolute_error(y_hat, y_test))
```

```
mae:  3.516311812818896e-14
```

```
In [17]: plt.scatter(y_test, y_hat)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel('Actual Charges')
plt.ylabel('Predicted Charges')
plt.title('Linear Regression - Actual vs. Predicted Charges')
plt.show()
```



```
In [18]: # To retrieve the intercept:
print(regressor.intercept_)
```

```
-8.526512829121202e-14
```

```
In [19]: # For retrieving the slope (coefficient of x):
print(regressor.coef_)
```

```
[ 3.68949741e-17 -1.39940559e-17 -1.49684473e-14  1.00000000e+00
  1.00000000e+00  1.00000000e+00  3.84079752e-16]
```

```
In [20]: y_pred = regressor.predict(X_test)
y_pred
```

```
Out[20]: array([ 66.    ,  51.    , 163.5   , ...,  63.75 , 236.625,  43.5   ])
```

```
In [22]: # To compare the actual output values for X_test with the predicted values
df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
df
```

Out[22]:

	Actual	Predicted
--	--------	-----------

110587	66.000	66.000
43371	51.000	51.000
84310	163.500	163.500
160055	212.250	212.250
172966	126.000	126.000
...
1248	116.625	116.625
48614	306.450	306.450
201198	63.750	63.750
194357	236.625	236.625
183327	43.500	43.500

41935 rows × 2 columns

```
In [21]: import statistics
SSE=sum((y_pred-y_test)**2)
SST = sum((statistics.mean(y_pred)-y_test)**2)
R2=1-SSE/SST
R2
```

Out[21]: 1.0